

Here's the receipt for one LLM call.

Itemised, so you can see what to cut.

AI CAFE • RECEIPT

Input tokens	10,000	\$0.0200
Output tokens	500	\$0.0050
TOTAL		\$0.0250





Fractions of a cent, **thousands a month**

\$0.025

per call

\$250

per day at
10,000 calls

\$7,500

per month (30
days)

**Small per-call savings turn into big monthly ones.
Let's find them.**



The only **formula** you need

```
cost = input × input price  
      + cached × cached price  
      + output × output price  
(prices are per 1M tokens)
```

Example rates in this deck: \$2 in, \$0.10 cached, \$10 out per 1M tokens (GPT-6.1 Sol standard).



One support-bot call, itemised

RECEIPT · 1 API CALL		
ITEM	TOKENS	COST
System prompt	2,000	\$0.0040
Chat history	1,800	\$0.0036
Retrieved docs	6,000	\$0.0120
User question	200	\$0.0004
Answer (output)	500	\$0.0050
TOTAL		\$0.0250

Where the **money** goes

Retrieved docs

48%

Answer

20%

System prompt

16%

Chat history

14%

User question

The user's question is under 2% of the bill. The context you add around it is the real cost.

Output tokens cost **5× more** here

1 output token
\$10 per 1M

=

5 input tokens
\$2 per 1M each

Long answers are expensive answers. Ask for the length you actually need.

Same call, 10,000 times a day

\$0.025

per call

↓ × 10,000 calls

\$250

per day

↓ × 30 days

\$7,500

per month

Now let's cut it, line by line.

Cache the **system prompt**

System prompt	2,000	\$0.0040
...cached instead	2,000	\$0.0002
SAVED / MONTH		\$1,140

Put stable text first so it gets cached. Reads cost 95% less; the first write costs a little more.

Retrieve fewer, better chunks

Retrieved docs	6,000	\$0.0120
...trimmed	3,000	\$0.0060
SAVED / MONTH		\$1,800

The biggest line on the receipt is the biggest lever. Rerank and keep only what helps.

Cap the answer length

Answer	500	\$0.0050
...capped	300	\$0.0030
SAVED / MONTH		\$600

Set a max output length, but only after checking that shorter answers still help users.





Not urgent? **Batch it.**

Standard

Real-time answers

PRICE

100%

Batch / Flex

Slower, background jobs

PRICE

50% lower

Reports, tagging and evals rarely need instant answers. Halve their cost.

The model choice is a **5×** decision

GPT-6 Astra
same call

\$0.125

GPT-6.1 Sol
same call

\$0.025

Use the cheapest model that passes your evals,
and route only the hard questions up.

Levers 1–3, applied

System prompt (cached)	2,000	\$0.0002
Chat history	1,800	\$0.0036
Retrieved docs	3,000	\$0.0060
User question	200	\$0.0004
Answer	300	\$0.0030
TOTAL		\$0.0132

\$7,500 → \$3,960 a month. That's 47% saved, with the same model.

3 things that **surprise your bill**

- 1 Very long prompts can cost more per token (e.g. over 272K input tokens on GPT-6.1 Sol)
- 2 Chat history grows every turn, and it's resent each time
- 3 Prices change, so re-check them before you budget

Estimate tokens at about 4 characters each, then measure the real counts in your logs.

OPINION

Print your receipt **before you optimise.**

Log tokens
per line item



Find the
biggest line



Cut that first

Most teams optimise the prompt wording. The money is usually in the context around it.

Cut your LLM bill: **cheat sheet**

Formula	tokens × price, for input, cached input & output
Cache	Stable prefix first: up to 95% off
Retrieve	Fewer, better chunks
Output	Cap the length; it costs 5× here
Batch	50% off non-urgent jobs
Model	The cheapest one that passes your evals

What's the biggest line on your receipt?

Context

docs & history

Output

long answers

Tell me in the comments.

 Follow for more

 Save this post

Vishal Gajula

 Repost